

Name: Dikshya Kafle

Student Number: 2018380039

Experiment 4:

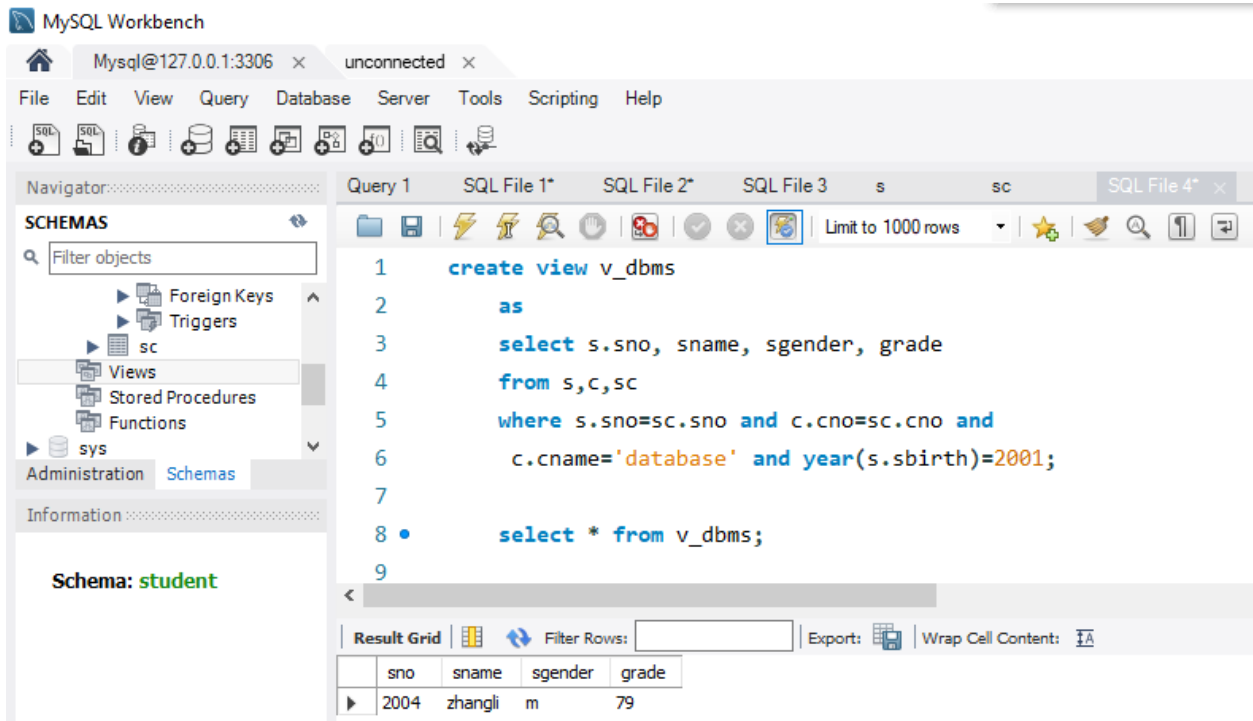
View and Index :

Goal

1. Be familiar with the use of GUI and SQL language to create, update and delete views.
2. Proficient in creating and deleting index using GUI and SQL language.
3. Understand and verify the role of index.

Content

1. In the student database, use SQL statement to create a view of students who have taken the database course and are born in 2001. The view includes the information of student number, name, gender and grade.



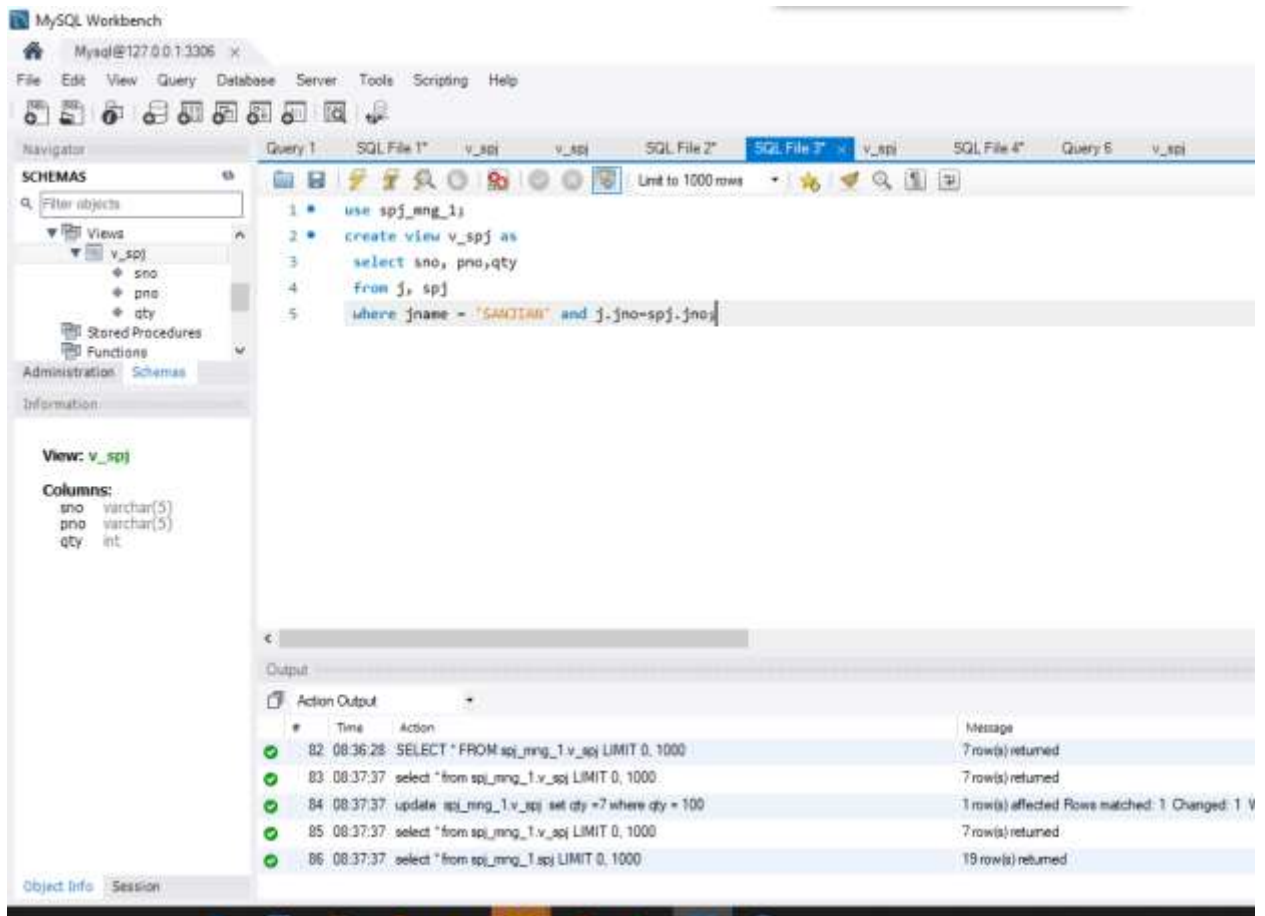
The screenshot displays the MySQL Workbench interface. The left sidebar shows the 'SCHEMAS' tree with 'student' selected. The main editor window shows the following SQL code:

```
1 create view v_dbms
2 as
3 select s.sno, sname, sgender, grade
4 from s,c,sc
5 where s.sno=sc.sno and c.cno=sc.cno and
6 c.cname='database' and year(s.sbirth)=2001;
7
8 select * from v_dbms;
9
```

The 'Result Grid' at the bottom shows the output of the query:

sno	sname	sgender	grade
2004	zhangli	m	79

2. Create a view for the supply situation of “SANJIAN” project, including the attributes of supplier code (SNO), part code (PNO) and supply quantity (QTY), with two ways with SQL statement.(view name:V_SPJ)



3. Complete the following view query with SQL statement.
(1) Find out all the parts code and their quantity used by “SANJIAN” project

MySQL Workbench

mysql@127.0.0.1:3306

File Edit View Query Database Server Tools Scripting Help

Navigator

SCHEMAS

Filter objects

Views

- v_spj
 - snq
 - pno
 - qty

Stored Procedures

Functions

Administration Schemas

Information

View: v_spj

Columns:

- snq varchar(5)
- pno varchar(5)
- qty int

Query 1

```

1 * Select pno,qty
2 from spj_eng_1.spj

```

Result Grid

pno	qty
P1	200
P1	100
P1	700
P2	100
P3	400
P3	200
P3	500

SQL 1

Output

Action Output

#	Time	Action	Message
83	08:37:37	select * from spj_eng_1.v_spj LIMIT 0, 1000	7 row(s) returned
84	08:37:37	update spj_eng_1.v_spj set qty = 7 where qty = 100	1 row(s) affected Rows matched: 1 Changed: 1 Warnings: 0
85	08:37:37	select * from spj_eng_1.v_spj LIMIT 0, 1000	7 row(s) returned
86	08:37:37	select * from spj_eng_1.spj LIMIT 0, 1000	19 row(s) returned
87	08:42:08	Select pno,qty from spj_eng_1.spj LIMIT 0, 1000	19 row(s) returned

(2) Find out the supply situation of supplier S1.

MySQL Workbench

mysql@127.0.0.1:3306 x

File Edit View Query Database Server Tools Scripting Help

Navigator

SCHEMAS

Filter objects

Views

v_spj

sno

pno

qty

Stored Procedures

Functions

Administration Schemas

Information

View: v_spj

Columns:

sno varchar(5)

pno varchar(5)

qty int

Query 1 SQL File 1* v_spj v_spj SQL File 2* SQL File 3* v_spj SQL File 4* Query 6 v_spj

1 • Select pno,qty

2 from spj_mng_1.spj

3 where sno='S1'

Result Grid

	pno	qty
P1	200	
P1	100	
P1	700	
P2	100	

spj 4 x

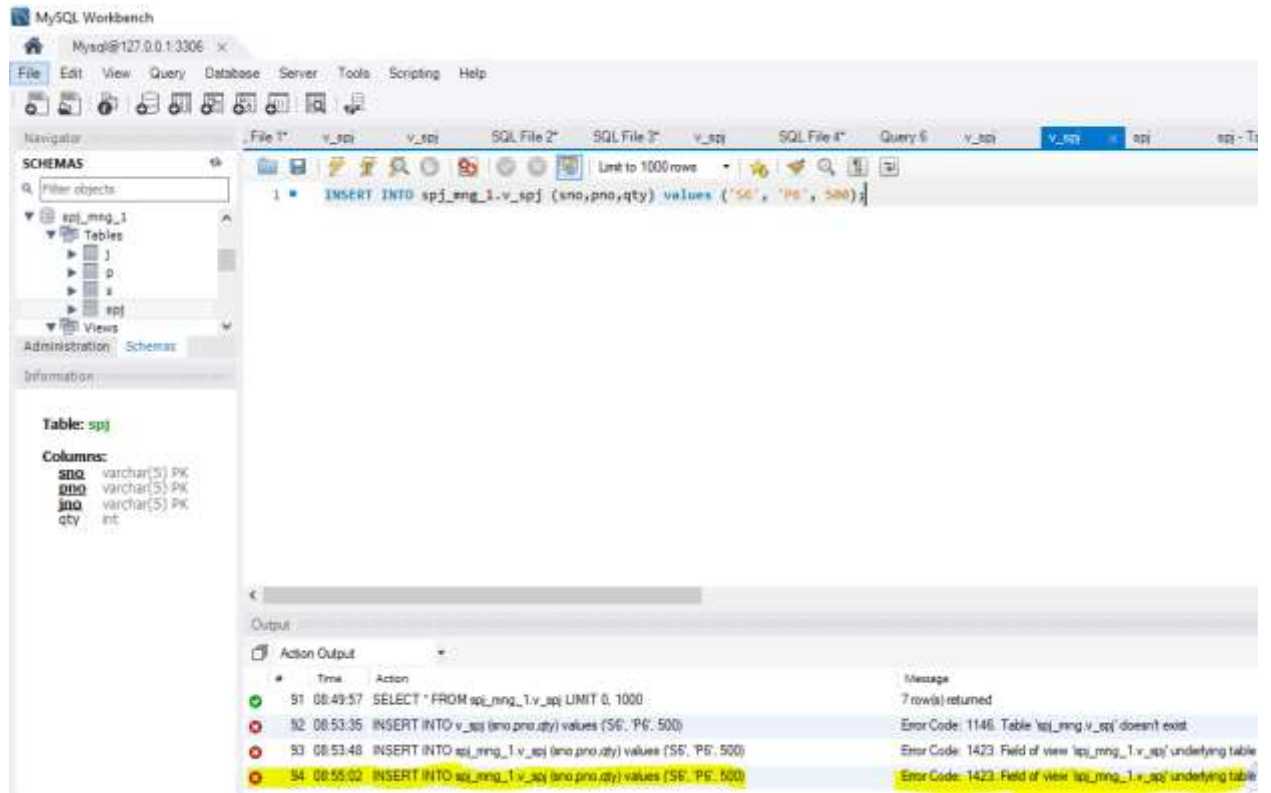
Output

Action Output

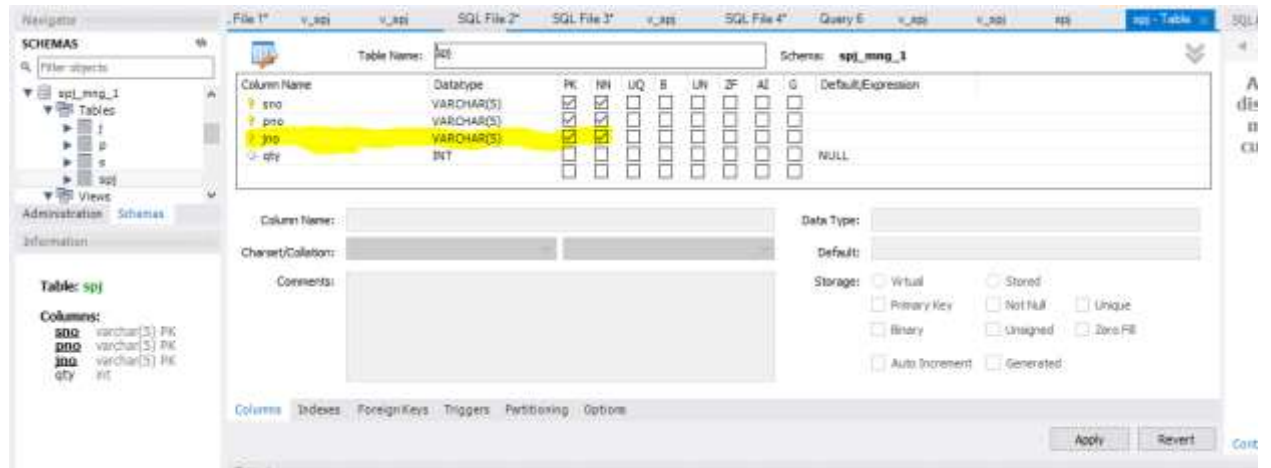
#	Time	Action	Message
86	08:37:37	select * from spj_mng_1.spj LIMIT 0, 1000	19 row(s) returned
87	08:42:08	Select pno,qty from spj_mng_1.spj LIMIT 0, 1000	19 row(s) returned
88	08:43:50	Select pno,qty from spj_mng_1.spj where sno='SN1' LIMIT 0, 1000	0 row(s) returned
89	08:44:49	Select pno,qty from spj_mng_1.v_spj where sno='SN1' LIMIT 0, 1000	0 row(s) returned
90	08:47:57	Select pno,qty from spj_mng_1.spj where sno='S1' LIMIT 0, 1000	4 row(s) returned

Object Info Session

4. Update the data of views with SQL statement.
 - (1) Insert a tuple into the view V_SPJ.



JNO is set to Not Null



- (2) Modify the quantity value of any tuple in the view V_SPJ.

MySQL Workbench

mysql@127.0.0.1:3306

File Edit View Query Database Server Tools Scripting Help

Navigator

SCHEMAS

Filter objects

Views

v_spj

sno pno qty

Stored Procedures

Functions

Administration Schemas

Information

View: v_spj

Columns:

sno varchar(5)

pno varchar(5)

qty int

Query 1

SQL File 1* v_spj v_spj SQL File 2* SQL File 3* v_spj SQL File 4* Query 6 v_spj

Limit to 1000 rows

```

1 * select * from spj_mng_1.v_spj;
2 * update spj_mng_1.v_spj set qty =7
3 * where qty = 100;
4 * select * from spj_mng_1.v_spj;
5 * select * from spj_mng_1.spj;
6

```

Result Grid

sno	pno	qty
S1	P1	200
S1	P1	100
S1	P1	700
S1	P2	100
S2	P3	400
S2	P3	200
S2	P3	500

Output

Action Output

#	Time	Action	Message
82	08:36:28	SELECT * FROM spj_mng_1.v_spj LIMIT 0, 1000	7 row(s) returned
83	08:37:37	select * from spj_mng_1.v_spj LIMIT 0, 1000	7 row(s) returned
84	08:37:37	update spj_mng_1.v_spj set qty =7 where qty = 100	1 row(s) affected Rows matched: 1 Changed: 1 Warnings: 0
85	08:37:37	select * from spj_mng_1.v_spj LIMIT 0, 1000	7 row(s) returned
86	08:37:37	select * from spj_mng_1.spj LIMIT 0, 1000	19 row(s) returned

Object Info Session

Result Grid

Filter Rows:

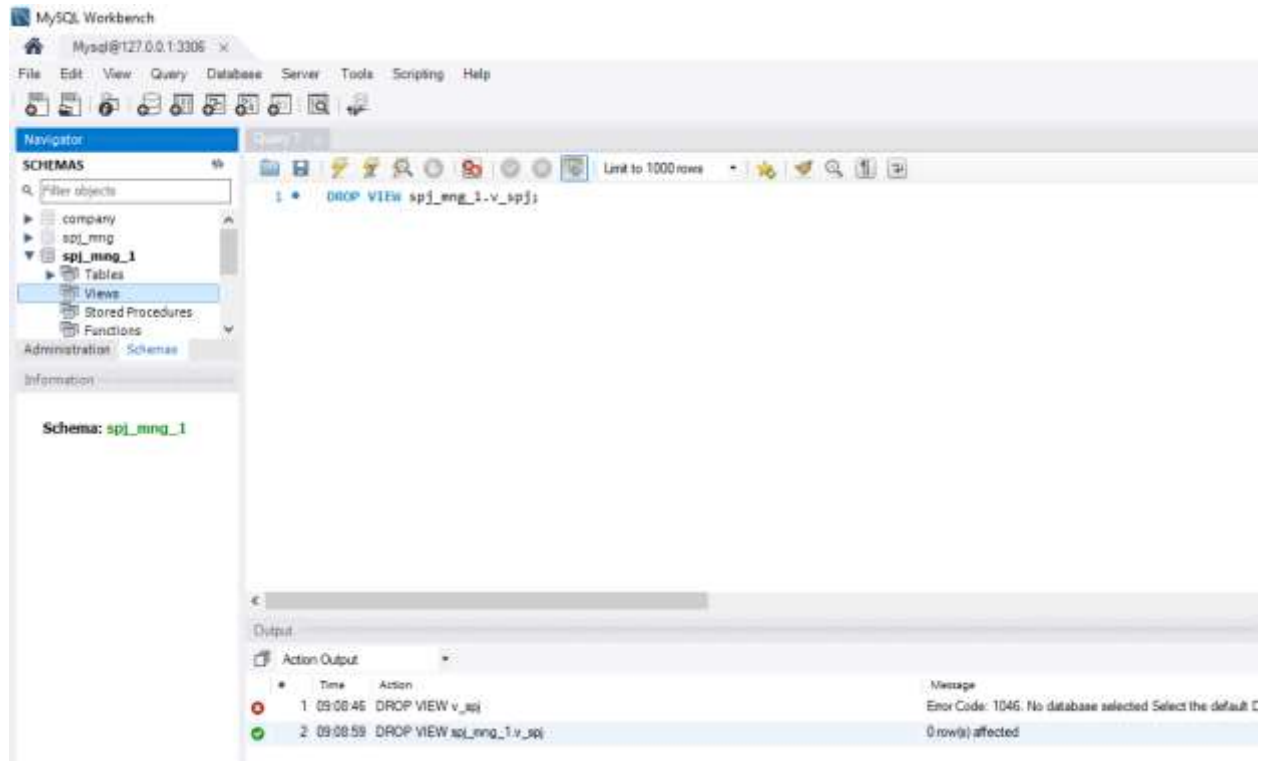
sno	pno	qty
S1	P1	200
S2	P3	400
S2	P5	400
S3	P1	200
S3	P3	200
S4	P5	7
S5	P3	200

v_spj 1

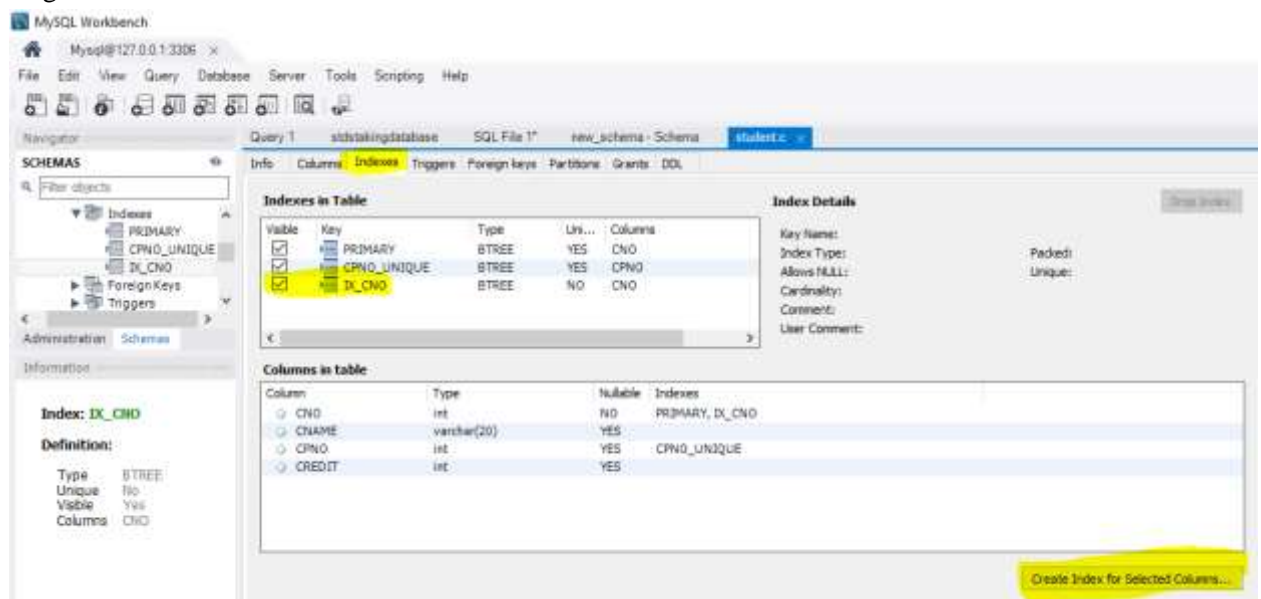
Output

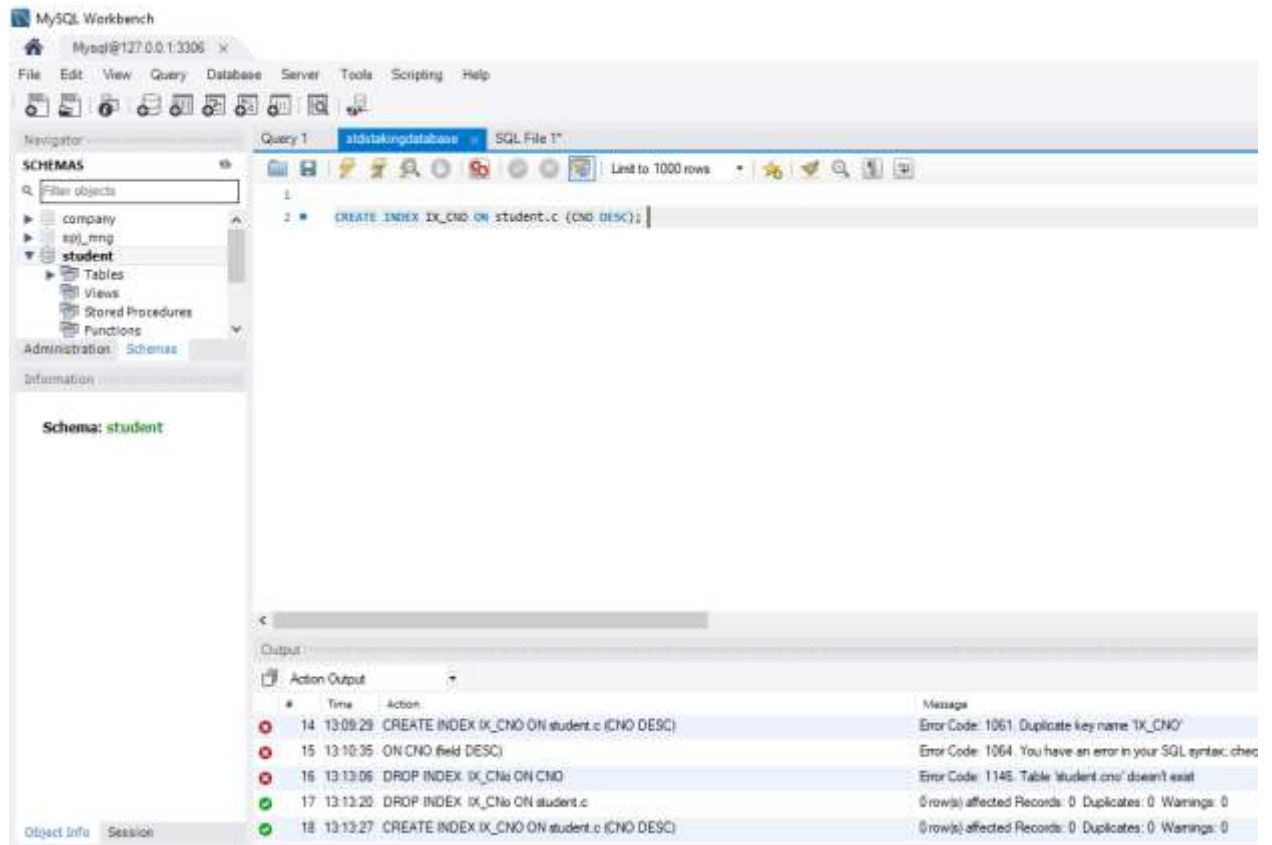
(3) Delete one tuple from the view V_SPJ

Hint: Only when the created view can be resolved, it can be deleted normally, otherwise it will fail to be deleted.

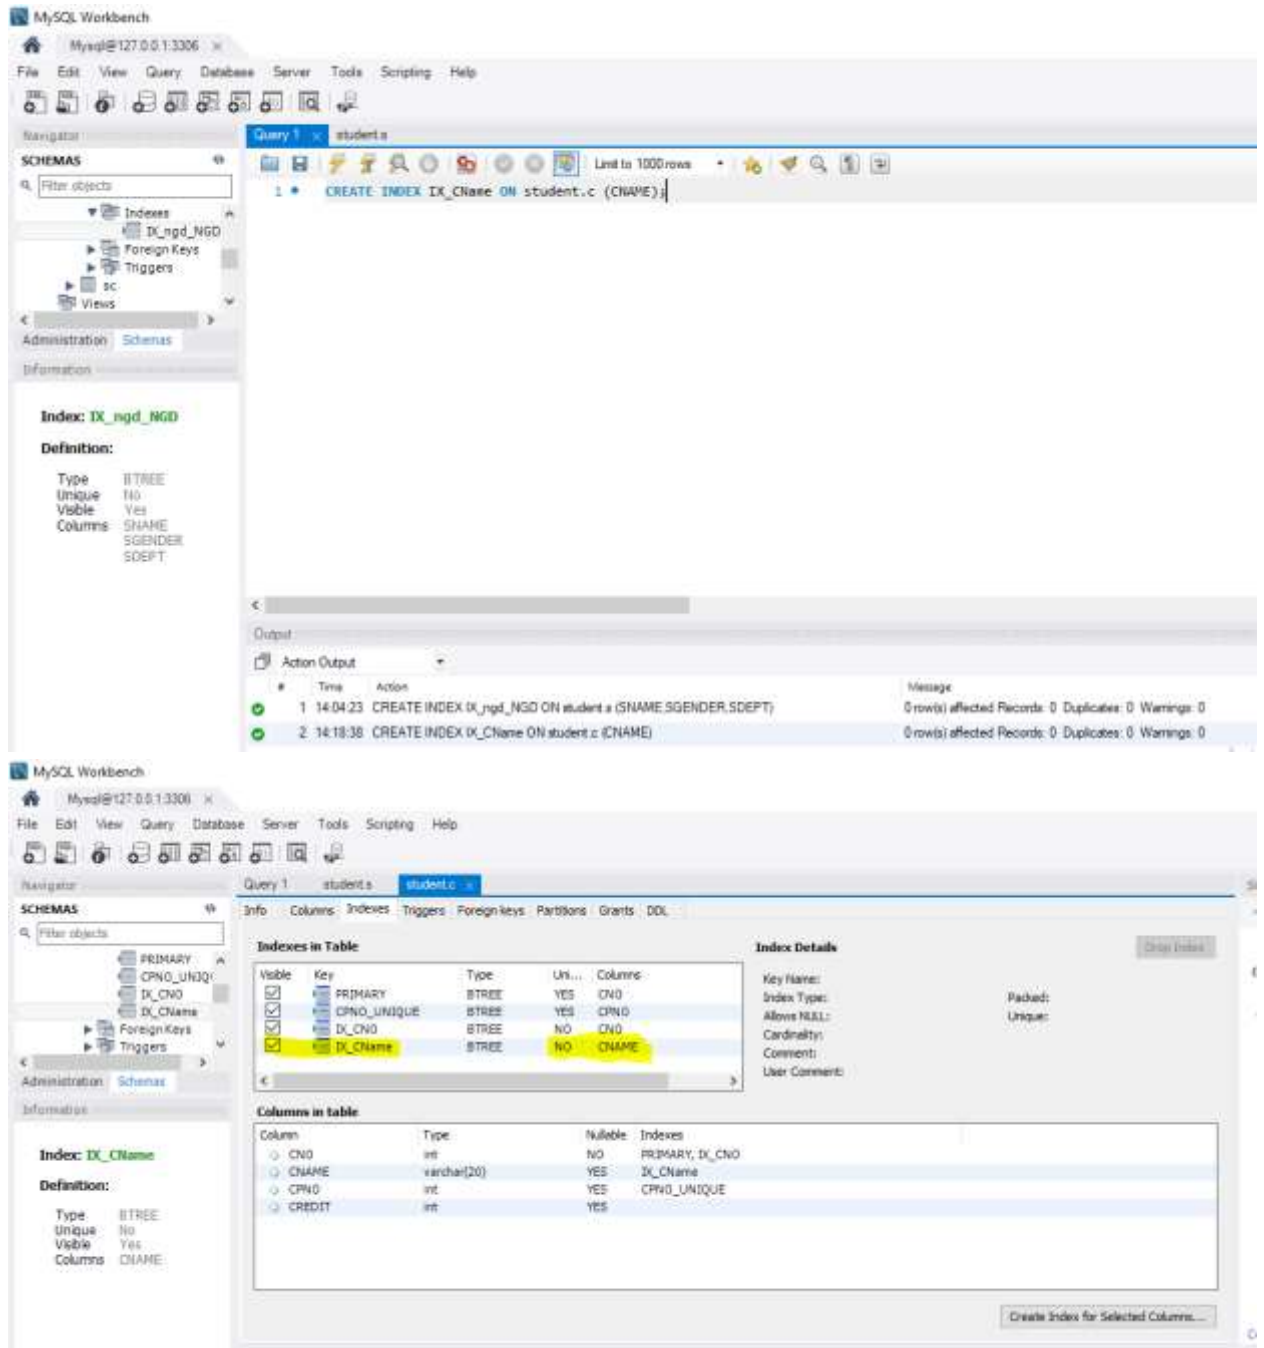


5. Create a descending index named IX_CNo for the CNO attribute of C table in student database by using GUI.

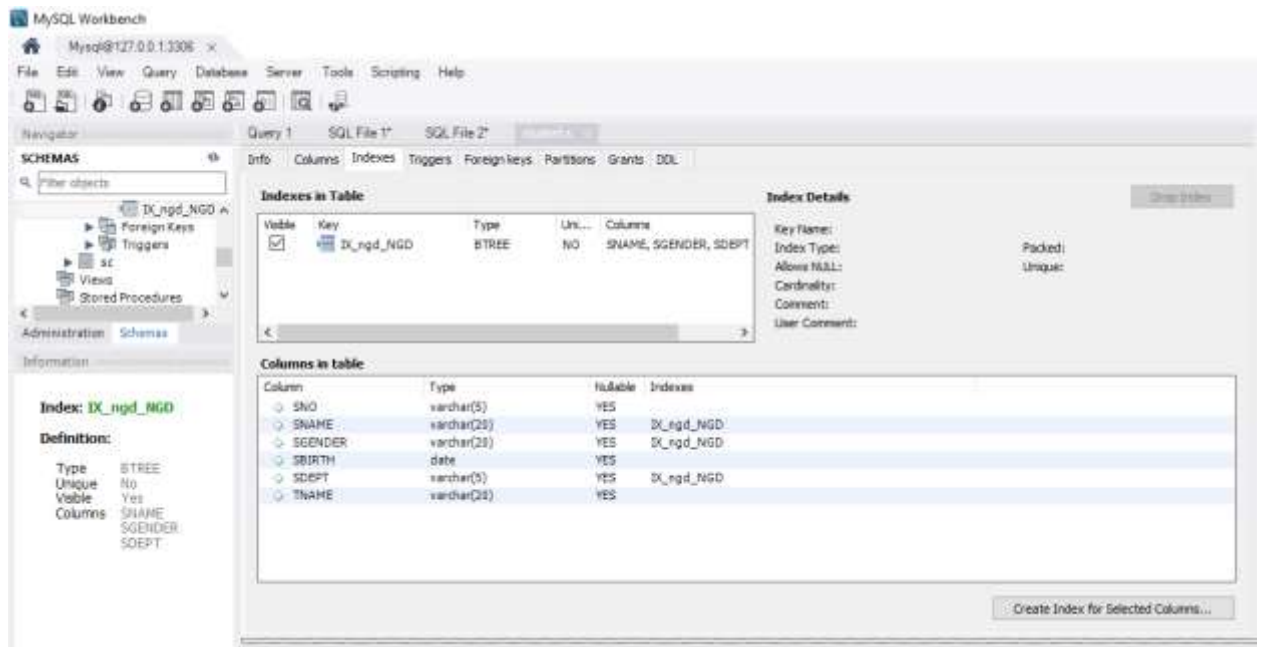
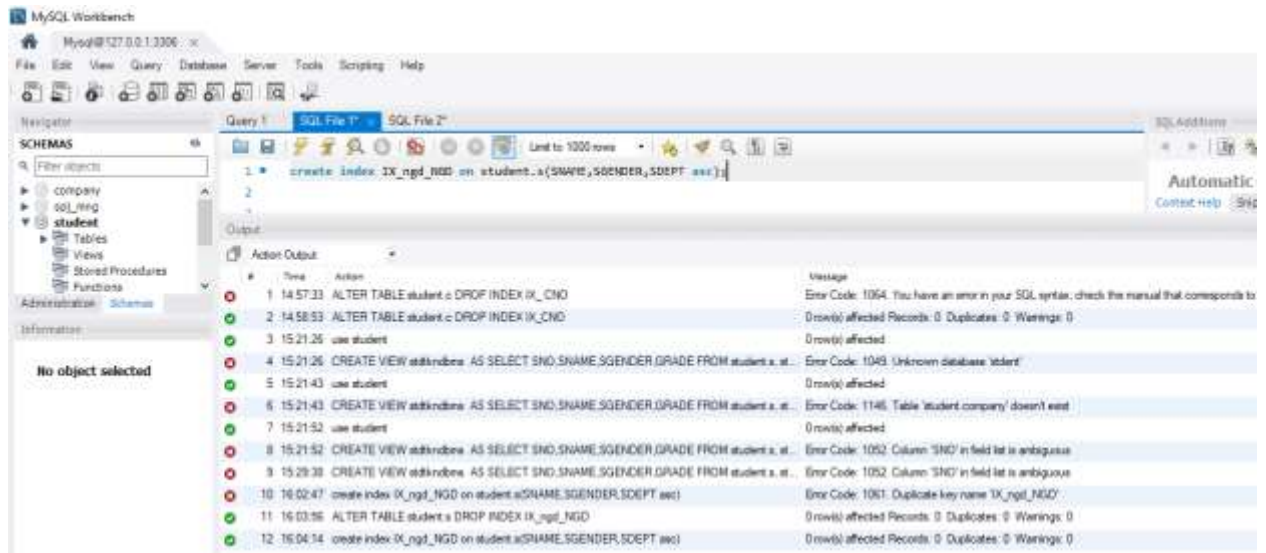




6. Use SQL statement to complete the following index operation on student database.
- (1) Create a non-unique index named IX_CName on the CNAME attribute of table C.



(2) Create a composite index named IX_ngd_NGD on the table S, which is an ascending index for sname, sgender and sdept attribute sets.



(2) Delete index IX_CNO of table C.



(4) Based on the above indexes (table C: primary key index of CNO, general index of CNAME; table S: primary key index of SnO, IX_ Nga composite index), use explain statement to obtain the query plan of each query statement, to observe the index usage in each query statement.

- ① explain select * from c;
- ② explain select * from c where cno = '1';
- ③ explain select * from c where cname='database' ;
- ④ explain select * from c where cname like '%database%';
- ⑤ explain select * from c where cname like 'database%';
- ⑥ explain select * from s where sname ='Zhangli' and sno='2001';
- ⑦ explain select * from s where sname ='Zhangli' and sgender='male' and sdept='IS';
- ⑧ explain select * from s where sname ='Zhangli' and sgender='male';
- ⑨ explain select * from s where sname ='Zhangli';
- ⑩ explain select * from s where sgender ='male';
- ⑪ explain select * from s where sgender ='male' and sdept='IS';

⑫

Navigation: Filter objects

- company
- spj_mng
- student
- sys
- tpc-h
- university
- user

Administration Schemas

Information

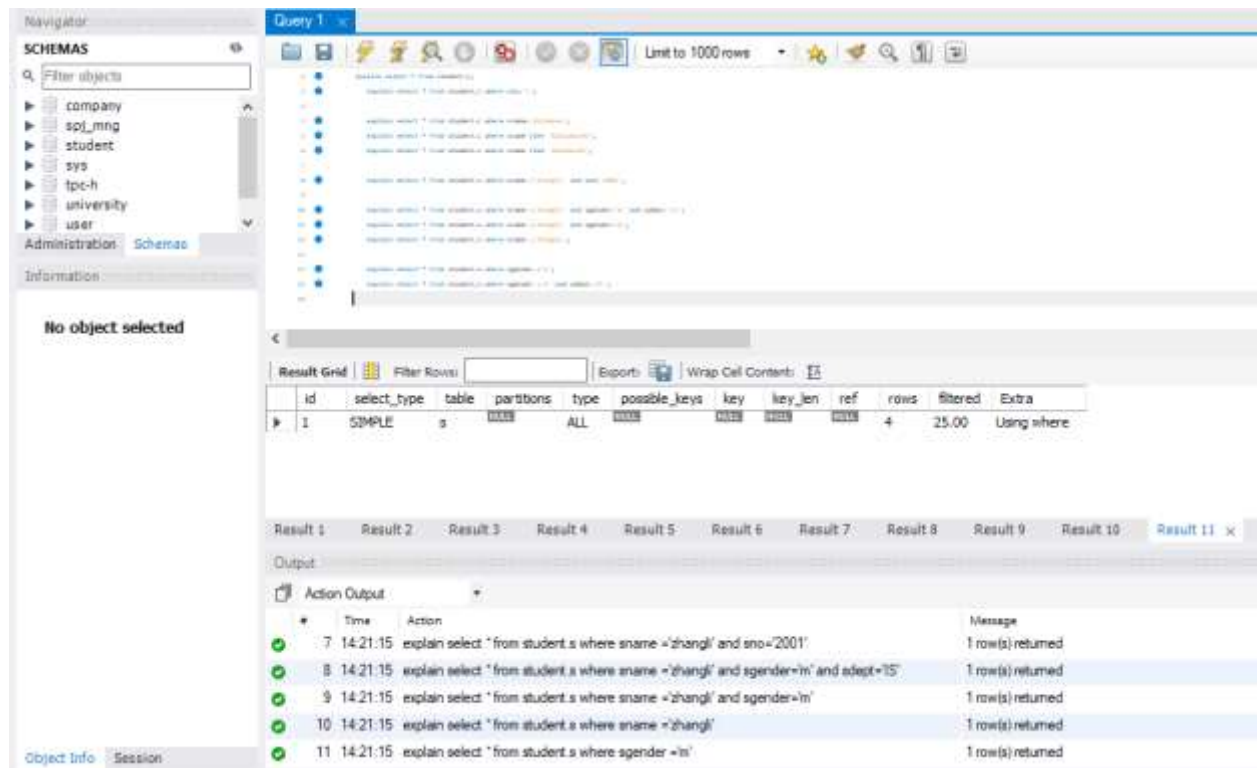
No object selected

Query 1

- 1 • explain select * from student.c;
- 2 • explain select * from student.c where cno='1';
- 3
- 4 • explain select * from student.c where cname='database';
- 5 • explain select * from student.c where cname like '%database%';
- 6 • explain select * from student.c where cname like 'database%';
- 7
- 8 • explain select * from student.s where sname ='zhangli' and sno='2001';
- 9
- 10 • explain select * from student.s where sname ='zhangli' and sgender='m' and sdept='IS';
- 11 • explain select * from student.s where sname ='zhangli' and sgender='m';
- 12 • explain select * from student.s where sname ='zhangli';
- 13
- 14 • explain select * from student.s where sgender ='m';
- 15 • explain select * from student.s where sgender ='m' and sdept='IS';
- 16

Result Grid Filter Rows: Export: Wrap Cell Content: [F4](#)

	id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
▶	1	SIMPLE	s	NULL	ALL	NULL	NULL	NULL	NULL	4	25.00	Using where



7. Suppose there is a basic table userinfo as follows, design an experiment to verify the effect of index on database query efficiency.

create table userinfo

(

user_id int primary key, //USER ID

username varchar(10), //USERNAME

gender char(1), //GENDER

age int, //AGE

c_id int //NO OF COLLEGE

)

Navigator: user - Schema user info table values

SCHEMAS

Filter objects

- tpc-h
- university
- user
 - Tables
 - userinfo
 - Views
 - Stored Procedures

Administration Schemas

Information

Table: userinfo

Columns:

user_id	int PK
username	varchar(10)
gender	char(1)
age	int
c_id	int

```

1 create table user.userinfo(
2   user_id INT PRIMARY KEY,
3   username VARCHAR(10),
4   gender CHAR(1),
5   age INT,
6   c_id INT
7 );
8 INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (1,'Ram','m',21,101);
9 INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (2,'Sita','f',23,102);
10 INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (3,'Hari','m',20,103);
11 INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (4,'Parbati','f',19,104);
12 INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (5,'Alix','f',22,105);
13 INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (6,'Sara','f',24,106);
14 INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (7,'Puna','m',22,107);
15 INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (8,'Kunti','f',19,108);
16

```

Output

Action Output

#	Time	Action	Message
9	19:07:46	INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (5,'Alix','f',22,105)	1 row(s) affected
10	19:07:46	INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (6,'Sara','f',24,106)	1 row(s) affected
11	19:07:46	INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (7,'Puna','m',22,107)	1 row(s) affected
12	19:07:46	INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (8,'Kunti','f',19,108)	1 row(s) affected
13	19:08:09	SELECT * FROM user.userinfo LIMIT 0, 1000	8 row(s) returned

Object Info Session

Result Grid

Filter Rows:

	user_id	username	gender	age	c_id
1	1	Ram	m	21	101
2	2	Sita	f	23	102
3	3	Hari	m	20	103
4	4	Parbati	f	19	104
5	5	Alix	f	22	105
6	6	Sara	f	24	106
7	7	Puna	m	22	107
8	8	Kunti	f	19	108

userinfo 1 x Apply

Output

Action Output

#	Time	Action	Message
9	19:07:46	INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (5,'Alix','f',22,105)	1 row(s) affected
10	19:07:46	INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (6,'Sara','f',24,106)	1 row(s) affected
11	19:07:46	INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (7,'Puna','m',22,107)	1 row(s) affected
12	19:07:46	INSERT INTO user.userinfo (user_id,username,gender,age,c_id) VALUES (8,'Kunti','f',19,108)	1 row(s) affected
13	19:08:09	SELECT * FROM user.userinfo LIMIT 0, 1000	8 row(s) returned

1. Verify the efficiency difference between indexed and non_indexed queries.

Indexed

The screenshot shows the MySQL Workbench interface. On the left, the 'SCHEMAS' pane shows the 'user' database with the 'userinfo' table. The 'Columns' pane for 'userinfo' lists: user_id (int PK), username (varchar(10)), gender (char(1)), age (int), and c_id (int). The main editor shows two SQL queries:

```
1 CREATE INDEX index_user
2 ON user.userinfo (username, gender, age, c_id);
```

The 'Output' pane shows the execution log:

#	Time	Action	Message
47	16:40:48	create view v_dbs_2 as select s.sno, s.name, s.gender, s.grade from s.c.ac where s.sno=s.c.sno	0 row(s) affected
48	16:40:52	select * from v_dbs LIMIT 0, 1000	0 row(s) returned
49	17:36:41	SELECT * FROM student.s LIMIT 0, 1000	4 row(s) returned
50	17:37:13	SELECT * FROM student.ac LIMIT 0, 1000	6 row(s) returned
51	17:41:14	SELECT * FROM student.ac LIMIT 0, 1000	6 row(s) returned
52	17:41:14	INSERT INTO student.ac (SNO,CNO,GRADE) VALUES (2004,1,79)	1 row(s) affected
53	17:44:11	drop view v_dbs	0 row(s) affected
54	17:44:16	drop view v_dbs_2	0 row(s) affected
55	17:47:12	create view v_dbs as select s.sno, s.name, s.gender, s.grade from s.c.ac where s.sno=s.c.sno	0 row(s) affected
56	17:47:12	select * from v_dbs LIMIT 0, 1000	1 row(s) returned
57	19:17:21	CREATE INDEX index_user ON userinfo (username, gender, age, c_id)	Error Code: 1146, Table 'student.userinfo' doesn't exist
58	19:17:49	CREATE INDEX index_user ON user.userinfo (username, gender, age, c_id)	0 row(s) affected Records: 0 Duplicates: 0 Warnings: 0

Drop index:

The screenshot shows the MySQL Workbench interface. The main editor shows the following SQL query:

```
1 DROP INDEX index_user ON user.userinfo;
```

The 'Output' pane shows the execution log, including the successful execution of the DROP INDEX statement:

#	Time	Action	Message
48	16:40:52	select * from v_dbs LIMIT 0, 1000	0 row(s) returned
49	17:36:41	SELECT * FROM student.s LIMIT 0, 1000	4 row(s) returned
50	17:37:13	SELECT * FROM student.ac LIMIT 0, 1000	6 row(s) returned
51	17:41:14	SELECT * FROM student.ac LIMIT 0, 1000	6 row(s) returned
52	17:41:14	INSERT INTO student.ac (SNO,CNO,GRADE) VALUES (2004,1,79)	1 row(s) affected
53	17:44:11	drop view v_dbs	0 row(s) affected
54	17:44:16	drop view v_dbs_2	0 row(s) affected
55	17:47:12	create view v_dbs as select s.sno, s.name, s.gender, s.grade from s.c.ac where s.sno=s.c.sno	0 row(s) affected
56	17:47:12	select * from v_dbs LIMIT 0, 1000	1 row(s) returned
57	19:17:21	CREATE INDEX index_user ON userinfo (username, gender, age, c_id)	Error Code: 1146, Table 'student.userinfo' doesn't exist
58	19:17:49	CREATE INDEX index_user ON user.userinfo (username, gender, age, c_id)	0 row(s) affected Records: 0 Duplicates: 0 Warnings: 0
59	19:20:02	DROP INDEX index_user ON user.userinfo	0 row(s) affected Records: 0 Duplicates: 0 Warnings: 0

• `EXPLAIN SELECT * FROM student.userinfo where username like 'm%';`

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	userinfo	NULL	ALL	NULL	NULL	NULL	NULL	6	16.67	Using where

Navigator

SCHEMAS

Filter objects

- company
- finance
- spj_mng_1
- student
- sys
- tpc-h
- university

Administration Schemas

Information

No object selected

Query 1

1 create index i_user on userinfo (username);

2 • Explain SELECT *FROM student.userinfo where username like 'm%';

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	userinfo	NULL	range	i_user	i_user	43	NULL	2	100.00	Using index condition

You can see the query execution efficiency that it searched for only 2 rows as shown in the figure after index formation, but before that it searched to fit the values in the table.

When we run the query, we check the CPU time (time spent on the query by CPU):

Note that elapsed time does not necessarily give the time of the query as the CPU might be busy with other processes and the query will be just waiting for the CPU so we just use the CPU time as the performance metric.

2. Verify the query efficiency of single field narrow index and multi field wide index, pay attention to understand the left most matching principle in wide index.

SQL Query:

To create three single field indexes on a table **test**:

```
SELECT * FROM student.test;
```

```
create index v1_test ON test(sr);
```

```
create index v2_test ON test(level);
```

```
create index v3_test ON test(value);
```

By using multiple operation on different indexes efficiency will be decreased because

DATABASE will fetch and manipulate data from 3 different views as compare to multi field index.

```
SELECT * from student.test where level='10' and value='233';
```

SQL Query:

To create one multi field wide index on a table **test**:

```
create index v123_test ON test(sr,level,value);
```

Testing:

It is more effective than the narrow single indexes after executing commands with 3 or 3 conditions that have 1 multi field index since it takes less time.

```
SELECT * from test where level='0' and value='234';
```

17:08:07.45 SELECT * from test where level='0' and value='234' LIMIT 0, 50000 5 row(s) returned 0.016 sec / 0.000 sec



3. Verify the difference of query efficiency between clustered index (primary key index) and secondary index: build clustered index and non-clustered index on the same field to compare query efficiency. (optional)

PRIMARY/CLUSTERED INDEX:



As we know by default in mysql PRIMARY index is clustered and all others are NON-CLUSTERED implicitly. Accessing a row using the clustered index is quick because it leads directly to the page with all data of row. If a table is large, it also helps often to save a disk I/O operation.

I have updated user ids but due to PRIMARY index its automatically updating the values to sort out the values of **user_id**:

Result Grid



Filter Rows:

Edit:

	user_id	username	gender	age	c_id
▶	1	Ram	m	21	101
	2	Sita	f	23	102
	3	Hari	m	20	103
	4	Parbati	f	19	104
	5	Alix	f	22	105
	6	Sara	f	24	106
	7	Puna	m	22	107
	8	Kunti	f	19	108

userinfo 1 ×

SECONDARY INDEX:

Many of the indexes are called secondary indexes, except the primary one. In order to search the data in rows, each secondary index contains the primary key column, we can pick a small primary key to reduce the execution time and improve execution efficiency.

4. At present, only memory engine of MySQL supports both b-tree index and hash index. Create a basic table based on memory storage engine in mysql, and verify the query efficiency difference between b-tree index and hash index based on this table. (optional)

B-TREE:

By default, when we create an index, its form is set to b-tree, which gives the best sort and helps to search for data such as binary search and others.

In the screenshot, you can see how to build the B-Tree Index and test the question after that:

SQL QUERY:

```
CREATE INDEX i_btree ON student.tab_innodb(sr);
```

Fetching data on base of range query help better and works efficiently with

The screenshot displays the MySQL Workbench interface. At the top, the SQL editor contains the query: `SELECT * from student.tab_innodb where industry like 'SFinance%';`. Below the editor, the 'Result Grid' shows a table with columns: sr, description, industry, level, size, line_code, and value. The data includes rows for 'Markets where business sold raw, unprocessed materials' and 'Markets where business sold manufactured or finished goods'. On the right, a message box states: 'Automatic context help disabled. Use the toolbar manually get help for the current caret position or toggle automatic help'. At the bottom, the 'Output' tab shows the 'Action Output' with three rows of execution logs, each indicating that 311 rows were returned in approximately 0.016 seconds.

sr	description	industry	level	size	line_code	value
35	Markets where business sold raw, unprocessed materials - ...	Finance	2	total	C0201.01	3
82	Markets where business sold raw, unprocessed materials - o...	Finance	2	total	C0201.02	0
129	Markets where business sold raw, unprocessed materials - ...	Finance	2	total	C0201.03	0
176	Markets where business sold raw, unprocessed materials - n...	Finance	2	total	C0201.04	198
223	Markets where business sold manufactured or finished good...	Finance	2	total	C0202.01	6
270	Markets where business sold manufactured or finished good...	Finance	2	total	C0202.02	0
317	Markets where business sold manufactured or finished good...	Finance	2	total	C0202.03	6

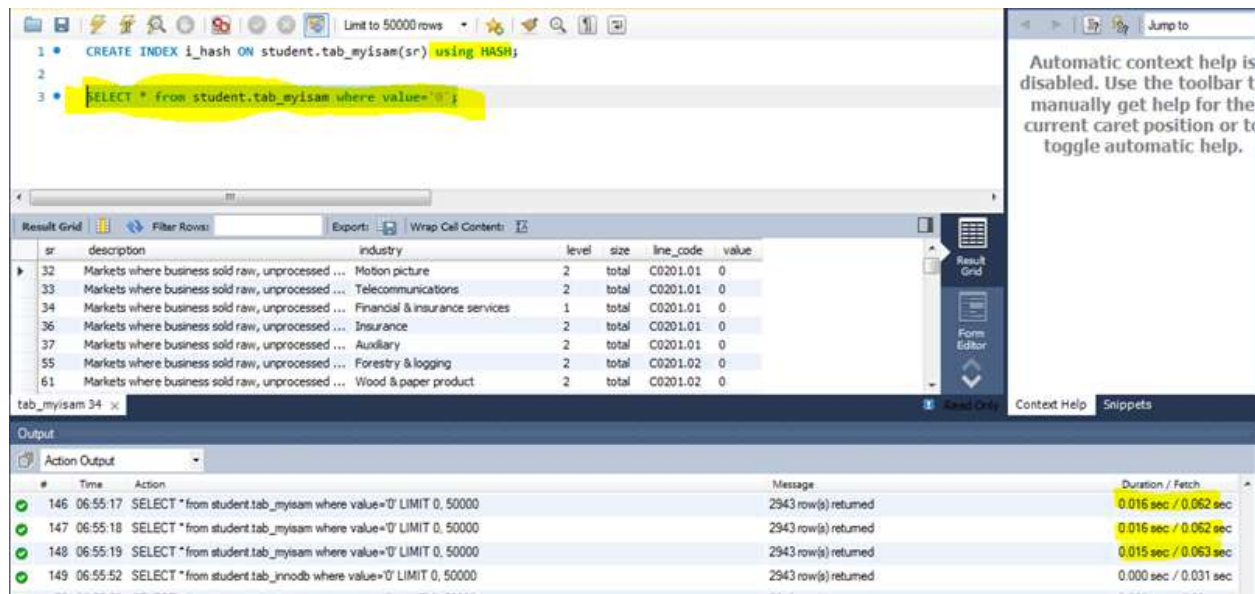
#	Time	Action	Message	Duration / Fetch
✓ 153	07:05:00	SELECT * from student.tab_innodb where industry like '%Finance%'; LIMIT 0, 50000	311 row(s) returned	0.000 sec / 0.015 sec
✓ 154	07:05:08	SELECT * from student.tab_innodb where industry like '%Finance%'; LIMIT 0, 50000	311 row(s) returned	0.016 sec / 0.016 sec
✓ 155	07:05:08	SELECT * from student.tab_innodb where industry like '%Finance%'; LIMIT 0, 50000	311 row(s) returned	0.000 sec / 0.016 sec

HASH INDEX:

Hash indexing is used for data fetching when equality operators are used, this indexing is best for such operations because they provide fast speed in data fetching.

SQL QUERY:

```
CREATE INDEX i_hash ON student.tab_myisam(sr) using HASH;
```



In general, we can say that for point queries, the hash index is best and the btree index works for range queries.

Conclusion:

We how the CPU time decreased when we added the index key to the appropriate attribute. This makes us able to conclude that indexes will make queries run faster and in fact, the difference is clearer as much as the table has a big number of rows (if we only deal with a small number of rows, adding an index will not have that big impact on the query performance)

Problem:

Especially, for the view and index experiment we faced a lot of technical failures. For instance, a simple code typing error led me to redo the whole database again. Also some of the questions were very challenging. I had problems in adding millions of data. At first I entered few data manually but it would take a lot of time. Adding an index does not result in reducing the cpu time.

Solution:

With, Prof. Xiaonan Zhao's very helpful guidance and assistance, I was able to solve most portions of our errors. I also used the internet and our textbook to find our remaining mistakes that led to fatal errors, it also provided us some SQL codes we haven't practiced.

Summary:

We practiced more on views and indexes. We created several views and tested how it is like to use the view by another user and checked the problems of updating a table through the view. After that, we practiced creating indexes with different properties using both GUI and SQL. Finally, we made an experiment to see how creating an index on an appropriate attribute can improve query performance in term of CPU time.